

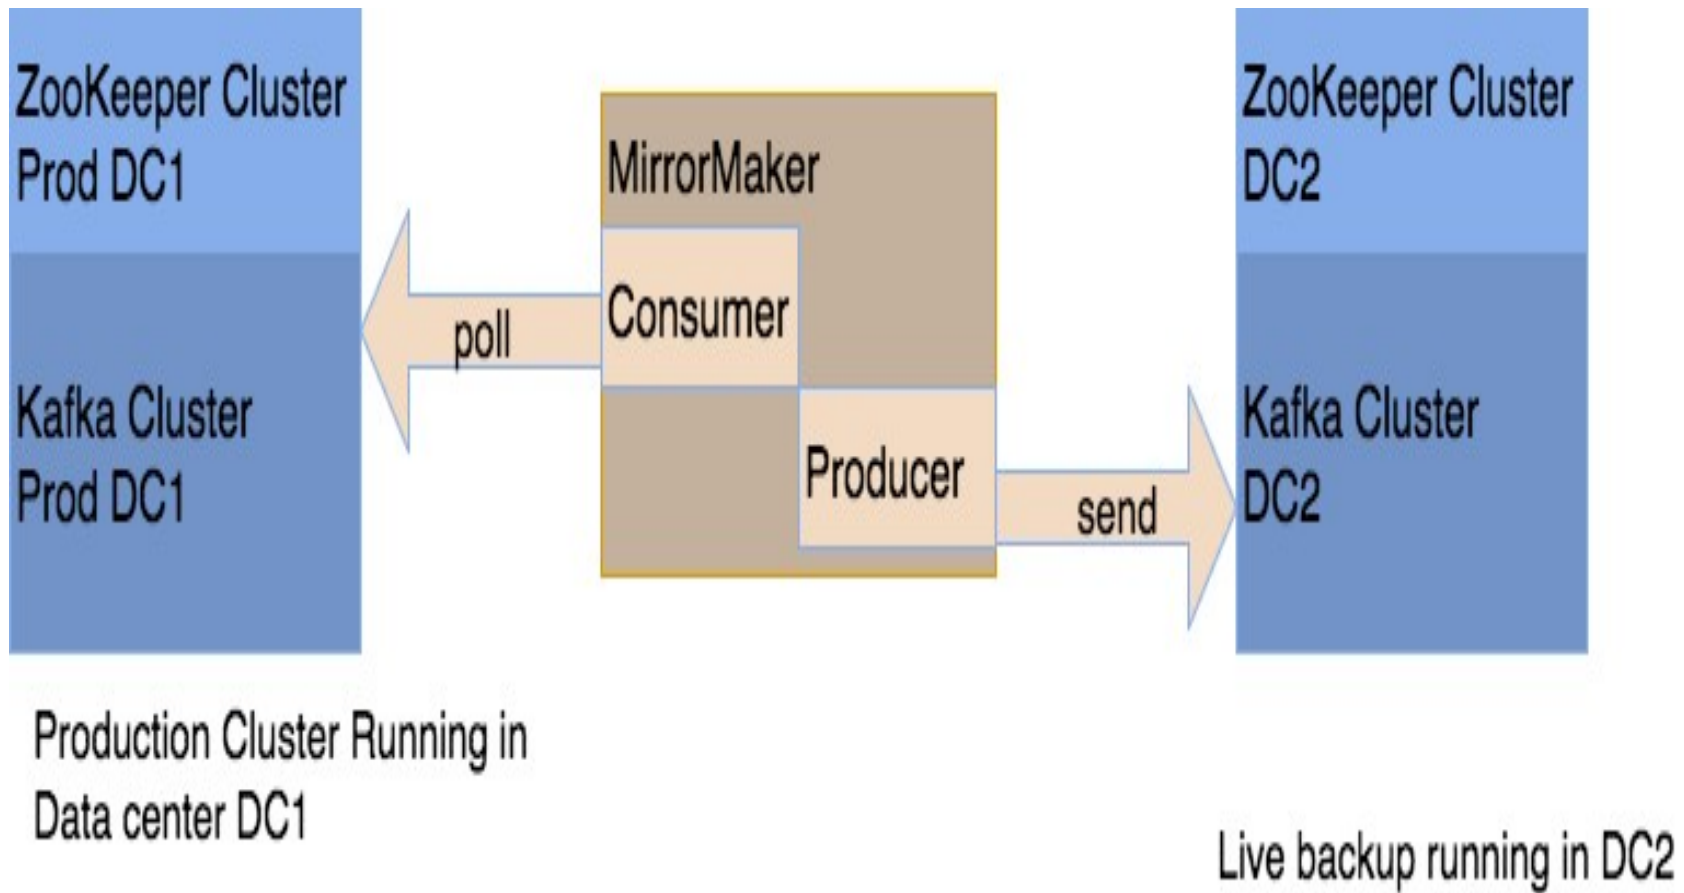
Kafka MirrorMaker

- ❖ Mirroring is replication between clusters – called ***mirroring*** to not confuse with ***replication***
- ❖ replication uses cluster involving brokers, partition leaders, partition followers, ISRs and ZooKeeper
- ❖ mirroring is just a consumer/producer pair in two clusters
- ❖ ***MirrorMaker*** is used for replicating one clusters data to another cluster

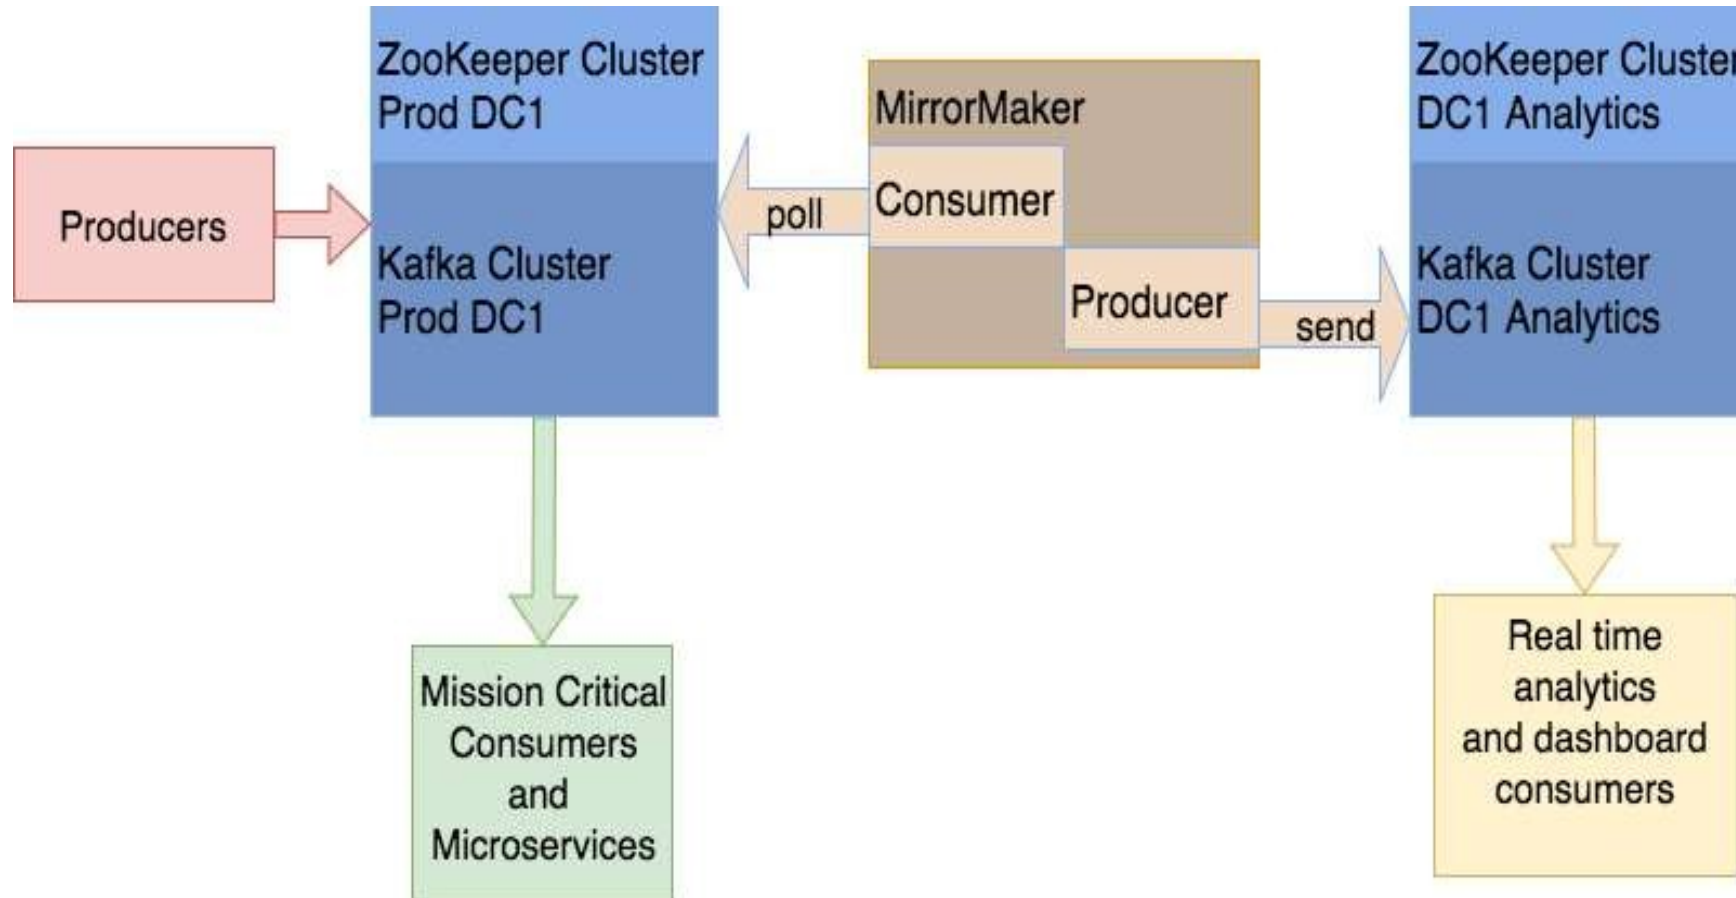
- Replicates topics (data plus configurations)
- Replicates consumer groups including offsets to migrate applications between clusters
- Replicates ACLs
- Preserves partitioning
- Automatically detects new topics and partitions
- Provides a wide range of metrics, such as end-to-end replication latency across multiple data centers/clusters
- Fault-tolerant and horizontally scalable operations

Disaster Recovery

Tos

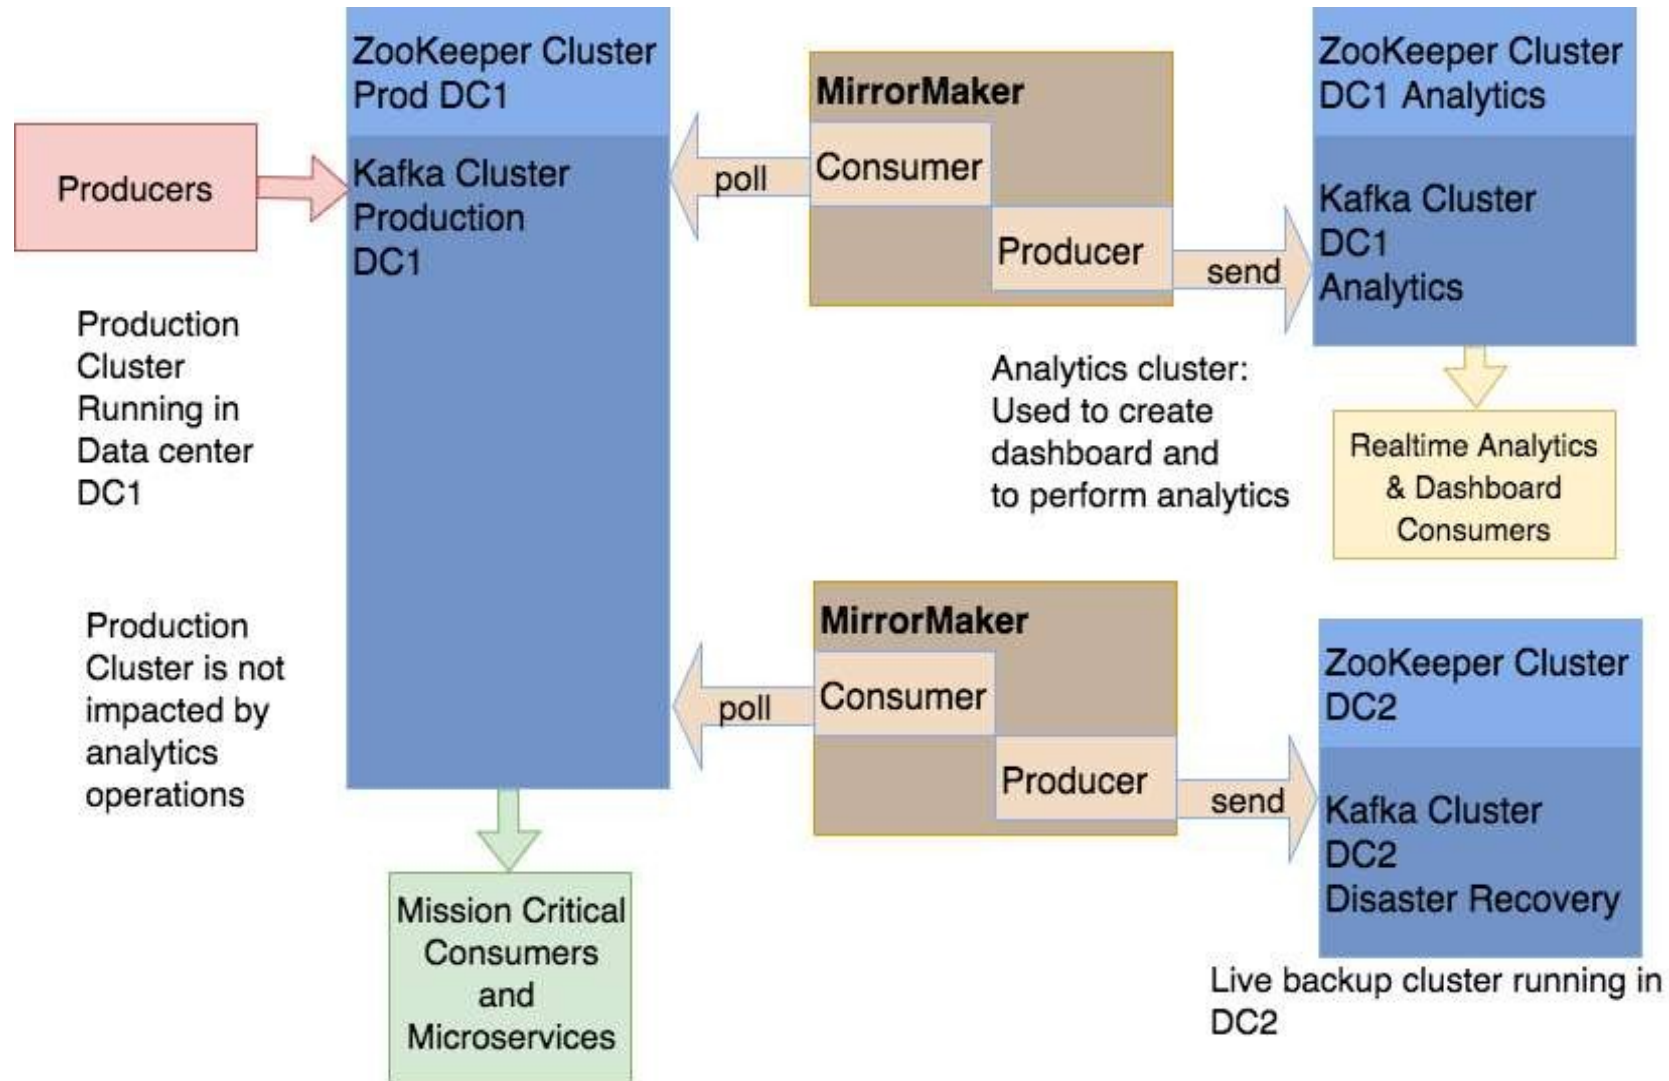


Scale Write, Avoid Impacting Mission Critical Services



Many MirrorMakers for different Purposes

Tos



```
1 | # Basic settings
2 | clusters = primary, secondary
3 | primary.bootstrap.servers = broker3-primary:9092
4 | secondary.bootstrap.servers = broker5-secondary:9092
5 |
6 | # Define replication flows
7 | primary->secondary.enabled = true
8 | primary->secondary.topics = foobar-topic, quux-.*
```

Run in secondary's data center, reading from the remote `primary` cluster

\$./bin/connect-mirror-maker.sh connect-mirror-maker.properties --clusters secondary